

DLA Draft

Developmental Learning Architecture: Learning History, Fast-Weight Traces, and Future Adaptation

First author: Yang Liu

Draft v0.2 (validation-stage) — workshop/arXiv

Abstract

We study whether a learner’s *prior learning history* changes how it will learn in the future, and which internal state carries this effect. Our Developmental Learning Architecture (DLA) keeps standard MLP/Transformer backbones unchanged but augments every weight with per-parameter fast weights, plasticity, and sleep consolidation, forming a fast/slow learner state. Using a 6.59M Chinese character-level GPT we show: (i) DLA matches tuned AdamW on new-domain adaptation while keeping consolidated slow memory essentially immune to forgetting; (ii) different curriculum histories (easy→hard vs hard→easy) lead to different future adaptation on an unseen domain; (iii) this history-dependent future adaptation is causally tied to the fast-weight state w_{fast} — replacing the hard→easy learner’s fast weights with those from an easy→hard learner significantly degrades future adaptation ($n=12$, gain@40 difference -0.019 , bootstrap CI excludes zero, Cohen’s $d \approx -0.7$); (iv) norm-matching and within-matrix shuffling do not remove the effect, so it is not explained by global magnitude or by a simple random-parameter control; (v) localization is suggestive of MLP and attention contributions; (vi) standard baselines (AdamW, replay, EWC) also show history

sensitivity, but DLA is the most robust after a difficult history; and (vii) the destructive w_{fast} effect replicates on a second Shakespeare character-level backbone. We do **not** find evidence that more experience monotonically accelerates learning: controlled longitudinal experiments across same-domain, cross-domain, and physics-near-transfer settings consistently fail to show a positive trend. The paper therefore proposes a precise framing: developmental change does not equal universally faster learning; rather, learning history leaves a persistent, partially localizable trace in transient learner state, with w_{fast} providing a causal component of history-dependent future adaptation.

1 Introduction

Standard neural training assumes the learner is a fixed function plus an external optimizer. This paper asks a more developmental question: *can the way a model learns change as a result of what it has learned?* We build a minimal architecture, DLA, in which each weight matrix is augmented by fast weights, plasticity, and sleep consolidation. The Transformer backbone itself is not modified.

Our empirical questions are:

- Does learning history change future adaptation on a new task?
- Which internal state carries that effect?
- Does the effect manifest as *faster learning* (monotonic self-improvement)?

We find support for the first two, and a clear negative result for the third.

2 Related Work

- Fast weights and fast/slow memory (Hinton & Plaut, 1987; Ba et al., 2016; Schmidhuber, 1992).
- Stability–plasticity dilemma and synaptic intelligence (Abraham & Robins, 2005; Zenke et al., 2017).
- Meta-learning and learned plasticity (Andrychowicz et al., 2016; Miconi et al., 2019; Beaulieu et al., 2020).
- Continual learning baselines used here: Elastic Weight Consolidation (Kirkpatrick et al., 2017) and experience replay (Robins, 1995).
- This work is distinct: we do not primarily propose a better continual-learning algorithm; we ask a mechanistic question about whether and

where “development” lives in the learner state.

3 Method

3.1 DLA state

For every 2D weight matrix (or embedding) we maintain:

- W_{slow} : consolidated long-term knowledge;
- W_{fast} : current learning trace;
- P : per-parameter plasticity gate;
- Q : slow-eligibility trace.

Effective weights:

$$W_{\text{eff}} = W_{\text{slow}} + \text{softplus}(P) * W_{\text{fast}}$$

The online (wake) update uses Adam-style moments on W_{fast} , modulated by $\text{softplus}(P)$, with P updated by a meta-plasticity rule. At task boundaries (sleep), W_{slow} is updated from Q and a small direct transfer of W_{fast} , and W_{fast} is decayed.

3.2 Histories and probe task

To create two different developmental histories we use three curriculum domains:

- easy→hard (EH): domains ordered from low birth perplexity to high;
- hard→easy (HE): the reverse order.

After experiencing all three tasks, the same model (stateful DLA) is probed on an **unseen domain D** under a fixed adaptation budget. We measure test perplexity trajectories and report gain@40 (relative PPL reduction after 40 steps), normalized learning efficiency, and T80.

4 Experimental Setup

Primary backbone: – 6.59M Chinese character GPT (6 layers, 8 heads, 256 dim, vocab 7280), pretrained on Chinese Wikipedia. – Curriculum domains: Wikipedia, SFT-style QA, science Wikipedia. – Unseen D: science slice not used in curriculum.

Second setting: – ~10.65M Shakespeare character GPT (6 layers, 6 heads, 384 dim, vocab 65).

Baselines: – AdamW; AdamW + replay; EWC, all using the same histories and D probe where applicable.

Metrics: – gain@40; LE_D; T80; initial PPL; per-task normalized loss slope; paired contrasts with bootstrap CI and sign tests.

5 Results

5.1 Fast/slow separation protects consolidated memory

Stage 4 formal (5 seeds): DLA matches tuned AdamW on adaptation to a new domain (gain +14.0% vs +13.0%) while consolidated slow-memory forgetting is ≈ 0 ($-0.4\% \pm 0.2$). This shows the fast/slow separation is a usable memory mechanism.

5.2 Learning history changes future adaptation

Across several settings, a hard→easy history leads to better future adaptation on unseen D than easy→hard. This history effect is not unique to DLA; standard learners also show it (Section 5.6), but DLA is the most robust after a difficult history.

5.3 Scalar learning-rule parameters are not a stable carrier

Body $\times$ ϕ cross-injection (Stage 5.5c): replacing the 9-dimensional tempo vector between histories changes future adaptation much less than replacing the body state. We therefore say ϕ is not a stable/dominant carrier in this setting, not that ϕ has no effect.

5.4 W_{fast} is a causal component

The core causal result (Stage 5.5e P0, n=12):

condition	mean gain@40
HE/HE	+0.014
HE + EH W_{fast}	-0.005

Paired contrast HE→raw EH: mean -0.019 , 95% CI $[-0.033, -0.005]$, Cohen's $d \approx -0.71$, sign $p=0.019$ (10/12 negative). Initial PPL differences are small; the effect appears in the adaptation trajectory.

5.5 Norm, shuffle and module controls

$n=12$ gain@40 paired contrasts (vs HE/HE):

control	mean Δ gain@40	95% CI	d
raw EH W_fast	-0.019	$[-0.034, -0.005]$	-0.71
norm-matched EH W_fast	-0.022	$[-0.034, -0.009]$	-0.90
shuffled EH W_fast	-0.017	$[-0.023, -0.009]$	-1.25
EH embedding W_fast	-0.006	$[-0.009, -0.002]$	-0.87
EH attention W_fast	-0.012	$[-0.018, -0.006]$	-1.06
EH MLP W_fast	-0.017	$[-0.025, -0.008]$	-1.05

Norm-matching does not rescue the effect; global magnitude is not the explanation. Shuffle also does not rescue the effect, so a simple random-parameter control is not sufficient. Module localization is suggestive of MLP and attention contributions; embedding contributes less.

Figure 7: W_fast cross-injection and full future-adaptation trajectories. A: norm/shuffle controls. B: module-localized transfer. C: paired gain@40 effects. D: example trajectories. Figure 7: W_fast cross-injection and full future-adaptation trajectories. A: norm/shuffle controls. B: module-localized transfer. C: paired gain@40 effects. D: example trajectories.

5.6 Fair baselines

$n=5$ (LE_D means):

method	EH	HE	HE-EH
AdamW	0.339	0.675	+0.336
AdamW + replay	0.375	0.961	+0.586
EWC	0.292	0.895	+0.603
DLA	0.517	0.709	+0.192

DLA is the most robust after an EH (easy→hard) history. Standard learners achieve higher HE performance but suffer larger EH degradation.

5.7 Second-setting replication

Shakespeare character GPT, 2 seeds:

seed	HE/HE gain40	HE + EH W_fast
0	+0.052	+0.004
1	+0.042	-0.001

The destructive W_fast effect replicates in direction on a different backbone/corpus. Small n; treated as exploratory support.

5.8 Negative result: development $\neq$ monotonically faster learning

We directly tested whether more experience makes later tasks easier to learn:

- Stage 5 (difficulty-normalized curriculum): LE roughly flat.
- Stage 6 (matched-difficulty science slices): 7/10 seeds directional, $p=0.17$.
- Stage 7 (cross-domain, 10 tasks, 20 seeds): regression slope ≈ 0 , $p=0.82$.
- Stage 8 (physics near-transfer, 20 seeds): $d=-0.17$.

We therefore do not claim that developmental learning produces universally faster learning. The evidence supports a more specific statement: learning history changes the learner's state and future adaptation, and this change is not simply "getting faster."

6 Discussion

- **What is safe to claim.** DLA separates fast adaptation from slow memory preservation. Learning history changes future adaptation. The effect is substantially body-mediated, with w_{fast} as a causal component. Norm and shuffle controls do not remove the effect. The effect is not monotonically faster learning.
- **What remains suggestive.** Exact localization to one module; the degree to which w_{fast} encodes content vs dynamics; generalizability beyond two small backbones.
- **What must not be claimed.** That w_{fast} is the only carrier; that ϕ has no effect; that development means universally faster learning; that

sleep consolidation strongly writes knowledge into w_{slow} ; that results at 6–10M parameters transfer to large language models.

- **Reviewer-facing limitations.** No pre-registration (we report $n=10 \rightarrow n=12$ extension transparently). Baselines $n=5$. Second setting $n=2$. One language plus Shakespeare English corpus only.
-

7 Conclusion

We provide evidence for the claim that a learner is not fully characterized by what it currently knows: its prior learning history can alter how it subsequently learns. The evidence points to a persistent, partially localizable trace in transient learner state, with w_{fast} providing a causal component of history-dependent future adaptation. We do not find evidence that this developmental effect simply manifests as faster learning.

Reproducibility

- Code: <https://github.com/JayCRL/DLA>
 - Experiment tables and result JSONs: see `docs/experiments.md`, `paper/overnight_report.md`, `paper/validation_audit.md`, `results/`.
 - Scripts:
 - `stage4_formal.py`, `stage5_progressive_curriculum.py`
 - `stage55e_wfast_p0.py`, `validation_p0_controls.py`, `validation_baselines.py`, `validation_second_setting.py`
 - `stage6_longitudinal.py`, `stage7_cross_domain.py`, `stage8_physics.py`
 - Figures: see `paper/figures/`.
-

References (to be completed)

- Hinton, G., & Plaut, D. (1987). Using fast weights to deblur old memories.
- Ba, J., Hinton, G., Mnih, V., Leibo, J., & Ionescu, C. (2016). Using fast weights to attend to the recent past.
- Zenke, F., Poole, B., & Ganguli, S. (2017). Continual learning through synaptic intelligence.
- Kirkpatrick, J., et al. (2017). Overcoming catastrophic forgetting in neural networks.

- Andrychowicz, M., et al. (2016). Learning to learn by gradient descent by gradient descent.
- Miconi, T., Clune, J., & Stanley, K. (2019). Differentiable plasticity.
- Beaulieu, S., et al. (2020). Learning to continually learn.
- Robins, A. (1995). Catastrophic forgetting, rehearsal and pseudorehearsal.