

Figure F -- continual-learning position (DLA swept over (fast_decay, lambda), not a single tuned point)

